

The Problem: Counterfactual F

An example:

The Problem: Counterfactual F

An example:

- Sensitive attributes influence screening and hiring

The Problem: Counterfactual F

An example:

- Sensitive attributes influence screening and hiring
- Algorithmic hiring systems offer opportunity to re

The Problem: Counterfactual F

An example:

- Sensitive attributes influence screening and hiring
- Algorithmic hiring systems offer opportunity to re
- **BUT:** they inherit societal biases from their training

The Problem: Counterfactual F


An example:

- Sensitive attributes influence screening and hiring
- Algorithmic hiring systems offer opportunity to re
- **BUT:** they inherit societal biases from their training

„How would a person’s CV look like if its sensit

What We Need: Causa

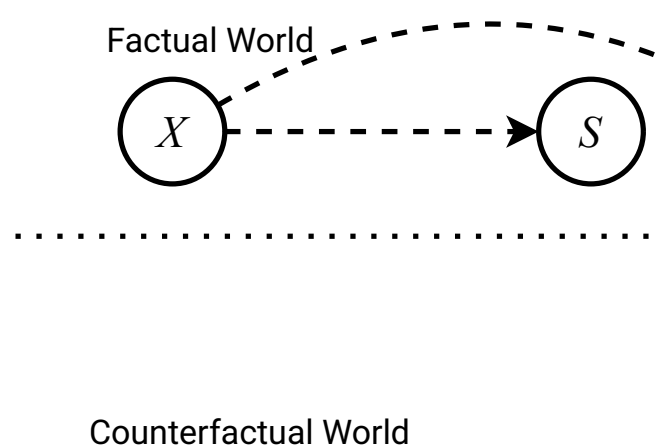
What We Need: Causa



Known:

X - text input (e.g., a CV)

What We Need: Causa

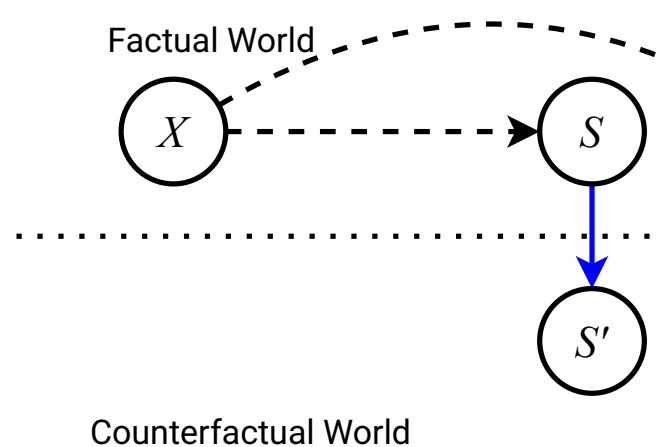


Known:

X - text input (e.g., a CV)

S - struct. features (e.g., extracted from CV)

What We Need: Causal



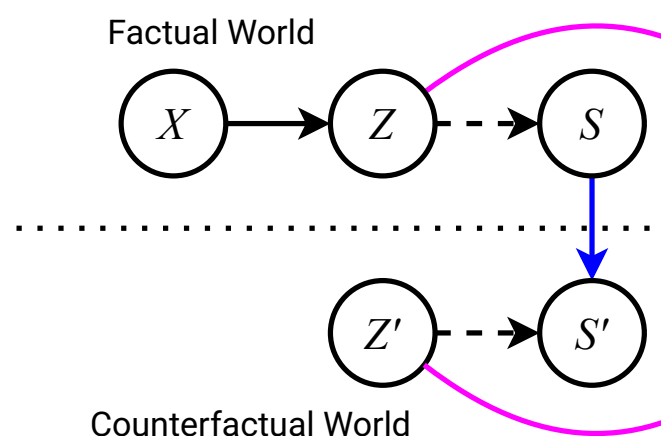
Known:

X - text input (e.g., a CV)

S - struct. features (e.g., extracted from CV)

h_S - counterfactual function in S space (**causal knowledge**)

What We Need: Causal



Known:

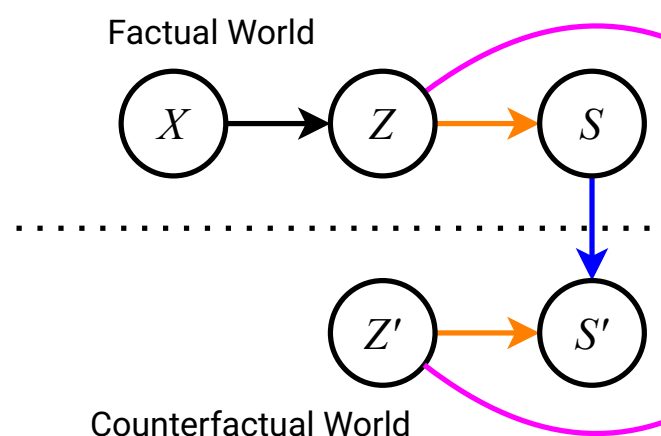
X - text input (e.g., a CV)

S - struct. features (e.g., extracted from CV)

h_S - counterfactual function in S space (**causal knowledge**)

$Z = f(X)$ - latent representation produced by a pre-trained encoder

What We Need: Causal



Known:

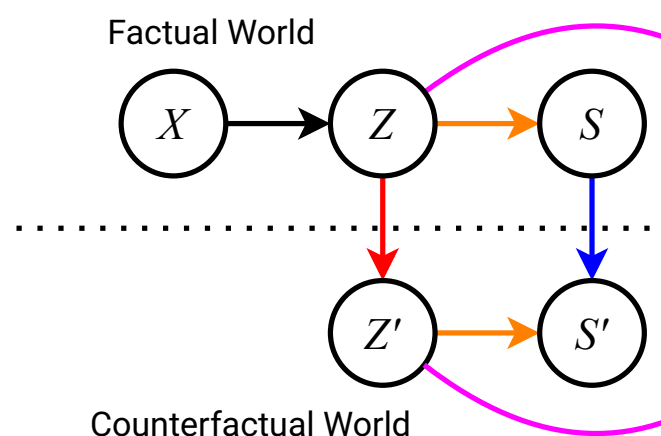
X - text input (e.g., a CV)

S - struct. features (e.g., extracted from CV)

h_S - counterfactual function in S space (**causal knowledge**)

$Z = f(X)$ - latent representation produced by a pre-trained encoder

What We Need: Causal



Known:

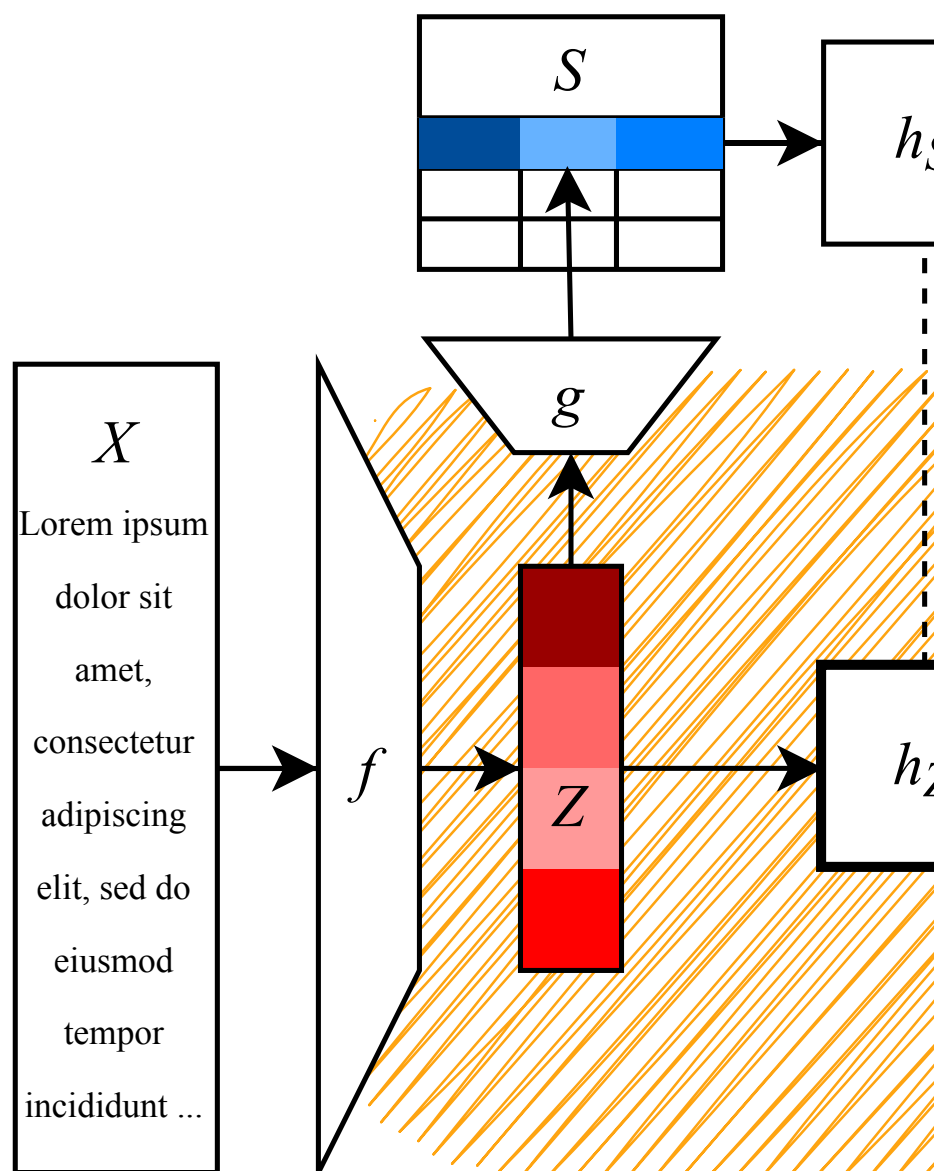
X - text input (e.g., a CV)

S - struct. features (e.g., extracted from CV)

h_S - counterfactual function in S space (**causal knowledge**)

$Z = f(X)$ - latent representation produced by a pre-trained encoder

The Idea: Modular Causally A



$$(h_S \circ g)(z) = (g \circ h_Z)(z)$$

Visit my poster to

- h_Z as a single-layer transformer vs. normalizing flow
- counterfactual point-estimates vs. distributions
- practicability of causal assumptions
- applications beyond counterfactual fairness